Hide menu

LSTMs and Named Entity Recognition

- **Video:** Week Introduction 1 min
- **Video:** RNNs and Vanishing Gradients 6 min
- **Reading:** RNNs and Vanishing Gradients 6 min
- **Reading:** (Optional) Intro to optimization in deep learning: Gradient Descent 10 min
- **Lab:** Vanishing Gradients 15 min
- **Video:** Introduction to LSTMs 4 min
- **Reading:** Introduction to LSTMs 3 min
- **Video:** LSTM Architecture 3 min
- **Reading:** LSTM Architecture 4 min
- **Video:** Introduction to Named Entity Recognition 3 min
- **Reading:** Introduction to Named Entity Recognition 2 min
- **Video:** Training NERs: Data Processing 4 min

Week 2 > Introduction to LSTMs

Previous Next

Introduction to LSTMs

The LSTM allows your model to remember and forget certain inputs. It consists of a cell state and a hidden state with three gates. The gates allow the gradients to flow unchanged. You can think of the three gates as follows:

Input gate: tells you how much information to input at any time point.

Forget gate: tells you how much information to forget at any time point.

Output gate: tells you how much information to pass over at any time point.

There are many applications you can use LSTMs for, such as:

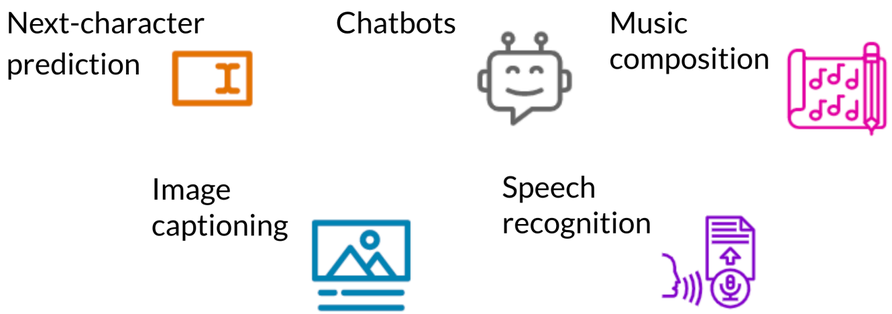

Here's a classic post on LSTMs with intuitive explanations and diagrams, to complement this week's material.

Mark as completed